

Implementation Pro — Product Rec

Version: 1.0

Last Updated: 2026-04-02

Owner: John Swapp

Status: Active — Living Document

AI agents that run implementations. You run the show.

Implementation Pro

Implementation Pro — Product Requirements Document (PRD)

Version: 1.0

Last Updated: 2026-04-02

Owner: John Swapp

Status: Active — Living Document

1. Executive Summary

Implementation Pro is a mobile-first, AI-powered SaaS platform that manages, governs, and executes SaaS implementations. It replaces manual spreadsheets, disconnected templates, and tribal knowledge with a unified system where AI agents handle repeatable billable work — and implementation teams oversee outcomes.

One-line pitch: AI agents that run implementations. You run the show.

Category: Implementation Management (distinct from Project Management)

2. Problem Statement

The Pain

SaaS implementation teams operate in a gap between sales hand-off and customer success hand-off. During that gap:

- No purpose-built tooling exists. Teams cobble together Jira, Smartsheet, Google Docs, and email. Nothing is designed for the implementation lifecycle specifically.
- Tribal knowledge walks out the door. When an experienced implementation consultant leaves, their playbook goes with them. New hires repeat the same mistakes.
- Scope creep is invisible until it's too late. Without a system enforcing scope boundaries, customers expand requirements informally through email and meetings. The team absorbs the cost.
- Status reporting is manual and time-consuming. Consultants spend hours each week compiling status updates, risk summaries, and client communications — billable hours lost to admin.
- Risk surfaces late. By the time a project is flagged red, the damage is done. Early signals (timeline drift, stakeholder disengagement, scope expansion) go undetected because nobody is watching the data.
- Operations lack guardrails. Resource over-allocation, unbilled time, and compliance gaps go unchecked until they become financial problems.

Who Feels This Pain

- Implementation Managers / Consultants running 3-8 concurrent customer engagements
- Professional Services Directors managing teams of 5-50 implementation specialists
- CS leaders inheriting projects from implementation without clean handoff documentation

What They Do Today

They use general-purpose PM tools (Asana, Monday, Jira), spreadsheets (MoSCoW trackers, RACI matrices, go-live checklists), email (status updates, stakeholder comms), and document editors (requirements docs, runbooks) — none of which talk to each other or learn from past engagements.

3. Product Vision

Build the operating system for SaaS implementation teams — a platform where:

- Every engagement follows a repeatable, governed process from kickoff to go-live
- AI agents handle the repeatable work (docs, comms, testing, analysis) while humans make the judgment calls
- The system gets smarter with every engagement (self-learning) and recovers from failures automatically (self-healing)
- Operations, governance, and execution are unified in a single workspace

Three Operational Layers

Layer	What It Does	Value
L1: Operations Automation	Enforces resourcing rules, time tracking policies, financial controls, compliance	Prevents operational blind spots before they become financial problems
L2: Delivery Governance	AI generates project plans, monitors delivery signals, surfaces risks early, sends client updates	Catches risk signals weeks before manual reviews would
L3: Work Execution	AI agents execute repeatable billable tasks (docs, configs, testing, migrations, comms, analysis)	Recovers 10-20 hours per engagement from manual admin work

Plus two cross-cutting capabilities:

- Self-Learning: Captures feedback on every AI output, builds org-specific context (writing style, risk patterns, client preferences), and refines agent prompts over time
- Self-Healing: Monitors system health, auto-retries failed operations, activates circuit breakers, and escalates anomalies

4. Target Users & Personas

Primary: Implementation Consultant ("The Operator")

- Runs 3-8 customer engagements simultaneously
- Needs to track scope, stakeholders, risks, and deliverables per engagement
- Spends 30-40% of time on admin (reports, docs, comms) instead of delivery
- Mobile access is critical — often at client sites or between meetings
- Values: efficiency, clear accountability, professional output

Secondary: Professional Services Director ("The Governor")

- Oversees a team of 5-50 consultants
- Needs cross-engagement visibility: resource utilization, financial health, risk posture
- Cares about compliance (time tracking policies, scope approval workflows, billing accuracy)
- Values: governance without micromanagement, early risk detection, team scalability

Tertiary: Customer Success Manager ("The Inheritor")

- Receives handoff documentation after go-live
- Needs clean transition: what was built, what decisions were made, what risks remain
- Values: complete handoff docs, relationship context, ongoing health visibility

5. Functional Requirements

5.1 Foundation (Core Platform)

ID	Requirement	Priority	Status
F-001	User authentication via Supabase Auth (email/password, OAuth)	Must	Scaffolded
F-002	Organization creation with slug-based routing	Must	Scaffolded
F-003	Organization member management with role-based access (owner, admin, member, viewer)	Must	Scaffolded
F-004	Engagement CRUD with status lifecycle (kickoff → in_progress → uat → go_live → complete)	Must	Scaffolded
F-005	Engagement health tracking (green/yellow/red + 0-100 composite score)	Must	Schema exists

ID	Requirement	Priority	Status
F-006	Scope item management with MoSCoW prioritization (must/should/could/won't)	Must	Scaffolded
F-007	Stakeholder registry per engagement (name, role, influence, communication preference)	Must	Scaffolded
F-008	Decision log per engagement (decision, context, impact, reversibility, status)	Must	Scaffolded
F-009	RACI matrix per engagement	Must	Scaffolded
F-010	Checklist system (kickoff + go-live) with task status, ownership, sign-off fields	Must	Scaffolded
F-011	Lessons learned capture with categorization (process, communication, technical, scope, timeline)	Must	Scaffolded
F-012	Status report generation (internal, customer-facing, executive)	Must	Scaffolded
F-013	Activity log for audit trail	Should	Schema exists
F-014	Automation rules (trigger + action)	Could	Schema exists
F-015	Mobile-first responsive UI across all views	Must	Not started

5.2 Layer 1: Operations Automation

ID	Requirement	Priority	Status
L1-001	Resource profiles (skills, capacity, rates, timezone)	Must	Schema exists
L1-002	Resource allocation per engagement (role, hours/week, date range)	Must	Schema exists
L1-003	Time entry tracking (date, hours, billable flag, category, description)	Must	Schema exists

ID	Requirement	Priority	Status
L1-004	Financial tracking per engagement (budget, consumed, margin target vs. actual)	Must	Schema exists
L1-005	Compliance rule engine (configurable rules, severity levels, actions: warn/block/notify/auto-correct)	Should	Schema exists
L1-006	Compliance violation tracking and resolution workflow	Should	Schema exists
L1-007	Support for billing models: fixed fee, T&M, milestone-based	Must	Schema exists

5.3 Layer 2: Delivery Governance

ID	Requirement	Priority	Status
L2-001	AI-generated project plans from engagement context (milestones, phases, assumptions, constraints)	Must	Schema exists
L2-002	Risk signal detection (timeline, scope, budget, stakeholder, technical, resource) with severity and confidence scoring	Must	Schema exists
L2-003	Automated client update drafting (weekly status, milestone reached, risk alert, go-live countdown)	Must	Schema exists
L2-004	Delivery signal monitoring (velocity, scope drift, stakeholder engagement) with trend analysis	Should	Schema exists
L2-005	Client update approval workflow (draft → approved → sent)	Must	Schema exists
L2-006	Health score computation from delivery signals	Should	Schema exists

5.4 Layer 3: Work Execution (AI Agents)

Six AI agents, each with a defined execution mode:

Agent	Type	Execution Mode	What It Does
Documentation Agent	documentation	Propose-and-wait	Generates requirements docs, meeting agendas, configuration runbooks, handoff documents
Communication Agent	communication	Propose-and-wait	Drafts status updates, escalation emails, scope change notifications, go-live announcements
Testing Agent	testing	Propose-and-wait	Generates test plans, test cases from scope items, requirements traceability matrices
Migration Agent	migration	Assist	Plans data migrations, field mappings, validation checklists, rollback procedures
Configuration Agent	configuration	Assist	Configuration checklists, as-configured docs, drift detection
Analysis Agent	analysis	Auto-execute	Cross-engagement analysis: patterns, utilization, scope creep trends, lessons synthesis

ID	Requirement	Priority	Status
L3-001	Agent definition registry with system prompts, I/O schemas, execution modes	Must	Schema + seed data exists
L3-002	Agent task queue (queued → running → completed/failed/cancelled/awaiting_approval)	Must	Schema exists
L3-003	Agent execution logging (tokens, model, duration, cost, output)	Must	Schema exists
L3-004	Agent artifact generation (document, test_report, migration_log, config_file, status_report, project_plan)	Must	Schema exists

ID	Requirement	Priority	Status
L3-005	Artifact approval workflow (draft → approved → delivered → archived)	Must	Schema exists
L3-006	Agent console UI for triggering, monitoring, and reviewing agent work	Should	Page scaffolded

5.5 Self-Learning

ID	Requirement	Priority	Status
SL-001	Feedback pipeline: accepted/modified/rejected on every AI output	Must	Schema exists
SL-002	Org-specific learning context storage (writing style, risk patterns, common issues, client preferences, process norms)	Must	Schema exists
SL-003	Outcome correlation tracking (signal → action → result)	Should	Schema exists
SL-004	Prompt refinement using accumulated feedback and context	Should	Not started

5.6 Self-Healing

ID	Requirement	Priority	Status
SH-001	System health checks (API availability, DB performance, agent success rate, queue depth, error rate)	Must	Schema exists
SH-002	Self-healing event tracking (auto-retry, circuit break, fallback, graceful degradation, auto-rollback, anomaly detection)	Must	Schema exists
SH-003	Error pattern recognition with auto-resolution suggestions	Should	Schema exists

6. Non-Functional Requirements

6.1 Performance

- Page load time: < 2 seconds on 4G mobile connection
- API response time: < 500ms for standard CRUD operations
- AI agent task completion: < 30 seconds for document generation
- Support for 1,000 concurrent users without degradation

6.2 Security

- All data encrypted at rest (Supabase default: AES-256)
- All data encrypted in transit (TLS 1.3)
- Row Level Security (RLS) on all tables — users can only access data within their organization
- API keys and secrets stored in environment variables, never in code
- Supabase Auth handles session management, token refresh, and OAuth flows

6.3 Scalability

- Multi-tenant architecture with org-level data isolation via RLS
- Stateless API design — horizontal scaling via Vercel serverless functions
- Database connection pooling via Supabase (PgBouncer)

6.4 Reliability

- Self-healing system monitors and auto-recovers from common failures
- Circuit breakers prevent cascading failures in AI agent pipeline
- Graceful degradation: if AI service is unavailable, manual workflows remain functional

6.5 Mobile-First Design

- All UI designed mobile-first, then scaled up for desktop
- Touch targets minimum 44px
- No horizontal scroll on viewports 375px and wider
- Offline-capable for viewing cached engagement data (future)

6.6 Accessibility

- WCAG 2.1 AA compliance target
- Keyboard navigable
- Screen reader compatible

7. Technical Architecture Summary

Component	Technology	Rationale
Frontend	Next.js 16 (App Router)	Server Components for performance, App Router for file-based routing
UI Framework	Tailwind CSS + shadcn/ui	Utility-first styling, consistent component library
Backend/DB	Supabase (PostgreSQL)	Auth, database, realtime, storage, edge functions in one platform
AI Engine	Anthropic Claude API (primary)	Best-in-class reasoning for document generation and analysis
AI Fallback	OpenAI (planned)	Redundancy for AI operations
Billing	Stripe	Subscription management, checkout, customer portal
Email	Resend	Transactional email delivery
Deployment	Vercel	Serverless hosting, edge network, preview deployments
Testing	Playwright	Cross-browser E2E testing including mobile viewports

Database Schema

33 tables across 6 domains:

- **Foundation (12 tables):** profiles, organizations, org_members, engagements, scope_items, stakeholders, decisions, raci_items, checklist_items, lessons_learned, status_reports, activity_log, automation_rules
- **L1 Operations (6 tables):** resource_profiles, resource_allocations, time_entries, financial_tracking, compliance_rules, compliance_violations
- **L2 Governance (4 tables):** project_plans, risk_signals, client_updates, delivery_signals
- **L3 Agents (4 tables):** agent_definitions, agent_tasks, agent_executions, agent_artifacts
- **Self-Learning (3 tables):** learning_feedback, org_learning_context, outcome_correlations
- **Self-Healing (3 tables):** system_health_checks, self_healing_events, error_patterns

41 custom enum types define controlled vocabularies across the schema.

RLS is enabled on all tables with org-scoped access policies. A helper function `is_org_member()` validates membership for all org-scoped queries.

8. Pricing Model

All billing through Stripe. Gumroad is deprecated.

Tier	Monthly Price	Agent Tasks/Month	Target User
Free	\$0	—	Individual consultants exploring the platform
Pro	\$49/mo	50	Solo consultants or small teams (1-3 people)
Team	\$149/mo	500	Implementation teams (4-15 people)
Enterprise	Custom	Unlimited	Large professional services organizations

All paid plans include a 30-day free trial.

9. Success Metrics

Product Metrics

Metric	Target	Measurement
Monthly Active Engagements per org	≥ 3	Engagements with activity in trailing 30 days
Agent Task Adoption	≥ 60% of eligible tasks use agents within 90 days	Agent tasks / manual equivalent tasks
Time Saved per Engagement	10-20 hours over engagement lifecycle	Self-reported + agent task duration vs. manual estimate
Net Promoter Score	≥ 50	Quarterly survey

Business Metrics

Metric	Target	Timeline
Paid Users	100	Month 6 post-launch
MRR	\$10,000	Month 9 post-launch
Churn Rate	< 5% monthly	Steady state
Trial-to-Paid Conversion	≥ 15%	Steady state

10. User Flows

10.1 First-Time User Flow

- User visits landing page → signs up (email or OAuth)
- Onboarding wizard: create organization → invite team (optional) → create first engagement
- Dashboard shows engagement overview with health status
- User navigates into engagement detail → tabbed interface for scope, stakeholders, RACI, checklist, etc.

10.2 Engagement Lifecycle Flow

- **Kickoff:** Create engagement → populate stakeholders, scope items (MoSCoW), RACI → run kickoff checklist
- **In Progress:** Track time → monitor delivery signals → generate status reports → manage scope changes
- **UAT:** Run testing agent for test plans → track checklist completion → monitor risk signals
- **Go-Live:** Execute go-live checklist → generate go-live communication → capture lessons learned
- **Complete:** Generate handoff documentation → archive engagement → feed data into cross-engagement analysis

10.3 AI Agent Flow

- User selects an engagement → opens agent console
- Chooses agent type (documentation, communication, testing, etc.)
- Agent receives engagement context as input → generates output
- If propose-and-wait: user reviews, approves/modifies/rejects → feedback captured
- If assist: agent provides recommendations, user takes action
- If auto-execute: agent runs on schedule, logs results, surfaces anomalies

11. Assumptions & Dependencies

Assumptions

- Target users are already running SaaS implementations and have established processes (this is not a "learn to implement" product)
- Users have reliable internet access (mobile data or Wi-Fi) for real-time features
- Organizations have 3+ concurrent engagements to justify the platform investment
- AI model quality (Claude) will continue improving, reducing need for manual output editing

Dependencies

- Supabase for database, auth, and realtime — platform availability is a hard dependency
- Anthropic Claude API for all AI agent features — API outage degrades L3 entirely
- Stripe for billing — required for any paid tier functionality
- Vercel for deployment — required for production hosting

- Resend for transactional email — required for client update delivery

12. Out of Scope (Won't)

- **Project management features** (Gantt charts, sprint boards, backlog management) — this is not a PM tool
- **CRM functionality** — no lead/opportunity tracking, no sales pipeline
- **Custom integrations with specific SaaS products** (Salesforce, HubSpot, etc.) — future consideration only
- **White-label / OEM offering** — not in the 12-month roadmap
- **On-premise deployment** — SaaS only
- **Real-time collaboration** (Google Docs-style co-editing) — not in initial release

13. Risks & Mitigations

Risk	Likelihood	Impact	Mitigation
AI output quality insufficient for professional use	Medium	High	Propose-and-wait mode ensures human review; self-learning loop improves over time
Solo developer velocity limits feature delivery	High	Medium	Phased build plan; AI-assisted development; focus on foundation before layers
Market education required (new category)	Medium	High	Build-in-public content strategy; community building; free tier for adoption
Supabase or Anthropic API outage	Low	High	Self-healing system with circuit breakers; graceful degradation to manual workflows
Scope creep on the product itself	Medium	Medium	MoSCoW prioritization applied to own roadmap; strict phase gating

14. Release Plan

Phase	Timeframe	What Ships
1. Foundation	Weeks 1-6	Auth, orgs, engagement CRUD, dashboard, mobile-first UI
2. Templates	Weeks 7-12	MoSCoW scope tracker, RACI matrix, checklists, status reports, lessons learned
3. L1: Operations	Weeks 13-18	Resource management, time tracking, financial tracking, compliance engine
4. L2: Governance	Weeks 19-26	AI risk detection, health scoring, delivery signals, automated client updates
5. L3: Agents	Weeks 27-34	Six AI agents, artifact generation, agent console
6. Self-Learning	Weeks 27-34	Feedback pipeline, org learning context, prompt refinement
7. Self-Healing	Weeks 27-34	Health monitoring, auto-recovery, circuit breakers, error pattern recognition

15. Open Questions

- **Domain:** Is implementationpro.com secured, or should alternatives be evaluated?
- **Mobile app:** Is a native mobile app needed, or is a responsive PWA sufficient for mobile-first?
- **AI model costs:** At scale (1,000 orgs × 50 agent tasks/month), what is the projected Anthropic API cost? Does the \$49/mo Pro tier cover it?
- **Gumroad migration:** What happens to existing Gumroad subscribers? Migration path, communication plan, timeline for sunset.
- **Competitive landscape:** Are there emerging competitors in the "Implementation Management" category that need to be tracked?

This is a living document. Updated as decisions are made and requirements evolve.